



# Google Container Engine

Google Cloud Platform



Google Cloud Platform

# Agenda

1

Kubernetes and Container Engine

2

Container Cluster Resources

3

Interacting with Container Engine

4

Lab

# Kubernetes

- Sometimes shortened to 'k8s'
- [Open source](#) orchestration system
- Based on over 10 years experience at Google
- Built for a multi-cloud world:
  - Public
  - Private
  - Hybrid
- Modular design
- Can run almost anywhere



# Google Container Engine

- Based on **open source** [Kubernetes](#) orchestration system
- Orchestrate and **schedule** Docker containers
- Consumes **Compute Engine** instances and resources
- Uses a **declarative syntax** to manage applications
- **Decouple** operational and development concerns
- Manages and maintains
  - Logging
  - Health management
  - Monitoring
  - Scaling



# Agenda

1

Kubernetes and Container Engine

2

Container Cluster Resources

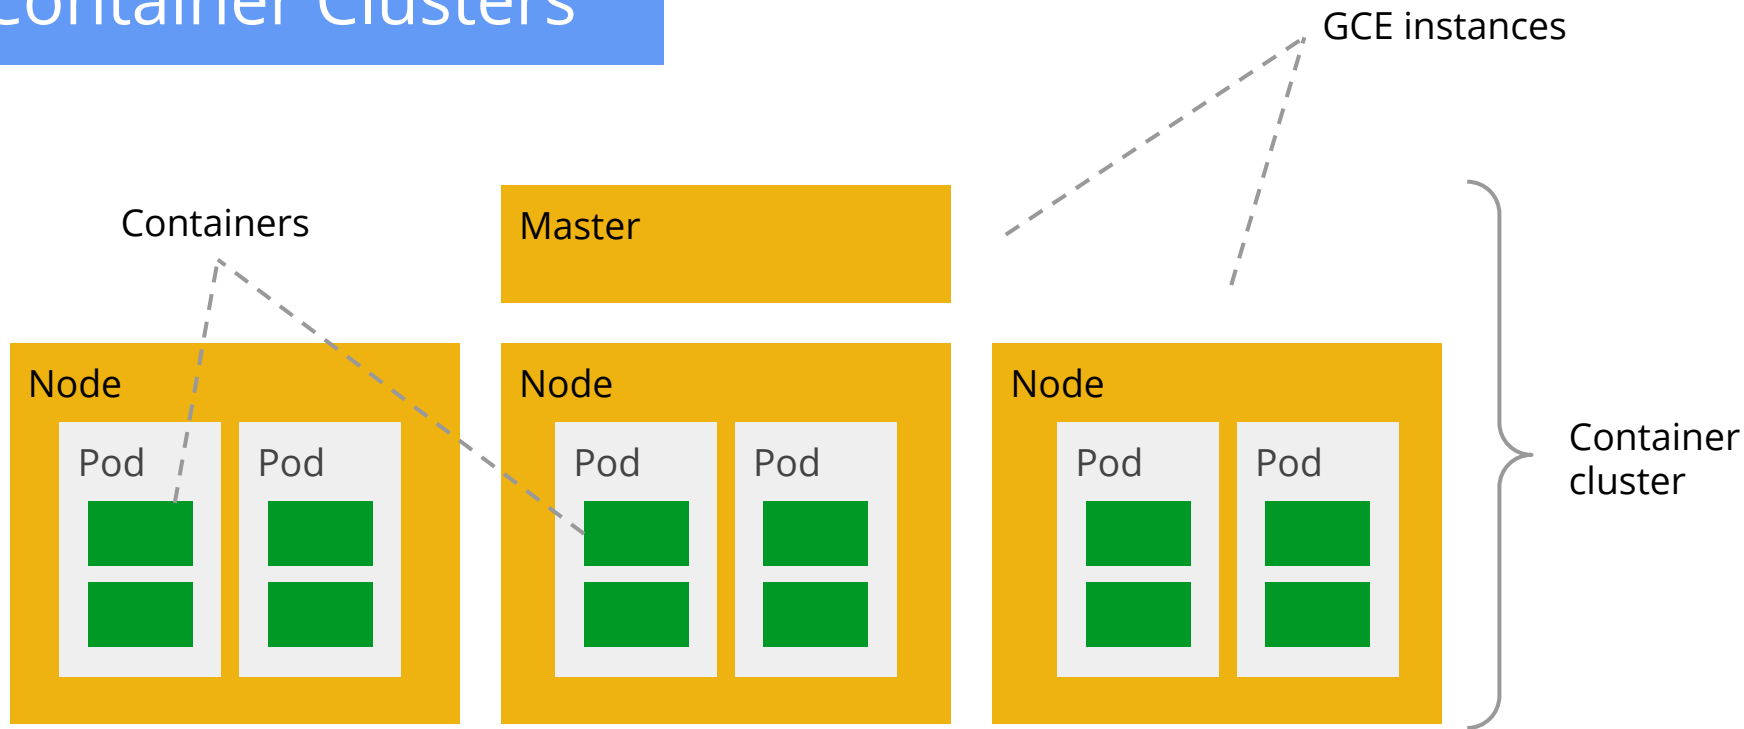
3

Interacting with Container Engine

4

Lab

# Container Clusters



Google Container Engine clusters are used to reliably scale container based workloads

# Cluster Resources



## Pods

Pods are **ephemeral** units that are used to manage one or more tightly coupled containers.

They enable **data sharing** and **communication** among their constituent components.



## Replication Controllers

Replication controllers create new pod "replicas" from a **template** and ensure that a configurable number of those pods are running.




## Services

Services provide a bridge based on an **IP and port pair** for non-Kubernetes-native client applications to access backends without needing to write code that is Kubernetes-specific.

# Pods


- When a host node instance dies, the pods scheduled to that node are deleted permanently
- Specific pods are never rescheduled to the new node instances
- Instead they must be replaced with new pods






# Replication Controllers

- Template settings are only used for new pod replica creation
- Changes to a template do not affect already-running pods or the switch to a new template
- Best practice to use a replication controller even for a single pod



# Services

- Pods are **ephemeral** and can be started or stopped at any time by the replication controller
- Pods will dynamically grow and shrink
- A service provides a stable **endpoint**



# Agenda

1

Kubernetes and Container Engine

2

Container Cluster Resources

3


Interacting with Container Engine

4

Lab

# kubectl

- kubectl is a command line utility that sends requests to the Kubernetes cluster manager
- Run **remotely** (not included in Cloud SDK)
- Frequently used kubectl commands include:  
get, describe, create, update, delete
  - Example - To get a list of pods:  
kubectl get pods



# Agenda

1

Kubernetes and Container Engine

2

Container Cluster Resources

3

Interacting with Container Engine


4

Lab

# Lab (1 of 2)

Deploy the Guestbook application on Container Engine.

1. Create a Container Engine cluster
2. Deploy a pod with the application containers
3. Test the application in your browser
4. Finish by deleting the container cluster



Images by Connie Zhou

# Lab (2 of 2)

GCP Project

GCE network

Master instance

Node

Pod

Docker

Frontend

Docker

Redis

Container Engine cluster comprised of 2 Compute Engine instances; a master & the node that will run the application pod

Will need to create firewall rule to allow traffic on port 80 (HTTP) again

# Resources

- Container Engine: Features Pricing Documentation  
<https://cloud.google.com/container-engine/>
- Kubernetes  
<http://kubernetes.io/>
- An introduction to containers, Kubernetes, and the trajectory of modern cloud computing  
<http://googlecloudplatform.blogspot.co.uk/2015/01/in-coming-weeks-we-will-be-publishing.html>
- Hello Wordpress  
<https://cloud.google.com/container-engine/docs/hello-wordpress>





cloud.google.com